

Solving the Airbus Ship Detection Challenge with U-Net

Armand Szokoly-Angyal, Lőrinc Toldi, Balázs Kiss

Abstract— Ship detection in satellite imagery is a key problem in maritime monitoring, naval security and global logistics. The Airbus Ship Detection dataset provides high-resolution images with run-length encoded segmentation masks for ships, posing challenges such as class imbalance, varying ship sizes and noisy annotations. In this project, we implemented a complete pipeline for ship localization and segmentation. Our primary solution uses a U-Net [1] architecture that performs semantic segmentation by predicting a binary ship mask for each image. To complement this, we built a lightweight convolutional neural network to classify whether an image contains ships at all, helping to mitigate data imbalance. We also experimented with a Mask Region Based Convolutional Neural Network (Mask R-CNN) [1] instance segmentation model to explore whether an object-level architecture could outperform classical semantic segmentation methods, although this approach ultimately performed poorly. Additionally, we conducted exploratory tests using the Segment Anything Model (SAM) [2] to compare automatic mask generation with ground-truth labels. The U-Net model achieved an F2 score [4] of 0.6 and became the final solution, while Mask R-CNN reached only 0.08, and SAM also produced negligible results. This documentation summarizes our dataset preparation, model architecture, training strategies, evaluation methods and the lessons learned from comparing multiple approaches.

I. INTRODUCTION

Accurately identifying ships in satellite images is a longstanding research topic with implications for commercial shipping oversight, illegal maritime activity detection and large-scale environmental observations. The Airbus Ship Detection dataset constitutes an attractive benchmark for such research, consisting of 768×768 RGB satellite images annotated with per-ship masks stored in run-length encoding format. The problem is challenging because the dataset contains many empty images, ships appear at vastly different scales and annotations vary in quality.

II. RELATED WORK

A wide range of approaches have been explored in prior work. Classical semantic segmentation architectures, particularly fully convolutional networks such as U-Net, have repeatedly demonstrated strong performance on dense prediction tasks, including remote sensing applications [3] [5]. At the same time, region-proposal-based architectures, especially Mask R-CNN [1], provide object wise predictions and could theoretically yield more structured outputs, a property that has been noted in maritime detection research as well [6]. Recent developments in foundation models, such as the Segment Any-

thing Model (SAM) [2], further broaden the landscape by offering general-purpose segmentation capabilities without task-specific training.

Our objective was to evaluate several practical pipelines for this dataset. We developed a U-Net model for semantic segmentation, a CNN classifier for detecting the presence of ships, and an experimental Mask R-CNN model for instance segmentation. These components allowed us to compare semantic and instance-aware approaches under the same processing conditions, supported by data validation and preprocessing.

III. THE MODEL ARCHITECTURE

Our primary model uses the U-Net architecture. It is a widely used CNN for image segmentation tasks, because it can capture both the overall structure and fine details of objects [3].

We resized all images and masks to 256×256 to reduce memory usage and speed up training and to emphasize more generic, important image features, even though using the full 768×768 resolution could produce more accurate segmentations. Similarly, we limited the dataset to a balanced sample of 10,000 images to save computational resources while maintaining a representative set of images with and without ships. Hyperparameters were chosen after testing multiple configurations to optimize segmentation performance.

The model outputs a single-channel probability map for each image, indicating the likelihood that each pixel belongs to a ship. To train the model, we use a combined loss function that equally weights Binary Crossentropy [7] and Dice loss (which is widely used in for instance medical segmentation tasks [8] to suppress class imbalance). Binary cross entropy encourages correct pixel-wise predictions, while dice loss directly measures the overlap between the predicted and true masks, which is particularly useful for handling the class imbalance between ship and background pixels. We also monitor Intersection over Union (IoU) [9] and F2 score as metrics, which provide a more holistic evaluation of segmentation quality by considering both precision and recall, rather than just pixel-wise accuracy.

The Adam optimizer [10] is used to update model parameters efficiently during training. Overall, this setup provides a balance between computational feasibility and segmentation performance, allowing us to experiment and evaluate the model within reasonable resource constraints.

Most of the images in this dataset are empty or mostly background. If we used simple accuracy as a metric, the model could just predict all zeros and still get a very high score.

Why not accuracy? Accuracy would mostly reflect background pixels and would make the model look better than it is.

IoU and F2 give a more realistic picture of how well the model is detecting ships.

As mentioned above, we implemented an instance-segmentation model based on Mask R-CNN with a ResNet-50 and FPN backbone, fine-tuned to detect ships as a single foreground category. Although this model's implementation, training loop and evaluation pipeline function correctly, its performance never reached practical levels, prompting us to treat it as an exploratory experiment.

IV. DATASET CLEANING

The dataset was downloaded using Google Drive mirrors, as Kaggle's API proved unreliable for multi-gigabyte transfers and hosting costs on commercial clouds were prohibitive. Once obtained, a detailed validation process ensured that the mask data met the assumptions of our models. Every RLE string was decoded and inspected for its number of connected components, its pixel area and whether it overlapped with other masks assigned to the same image. Masks that contained multiple disconnected shapes, failed to meet a minimum size threshold or overlapped with neighboring masks were corrected or discarded. This step significantly reduced annotation noise.

V. TRAINING

Training the U-Net required careful hyperparameter selection. A learning rate of 10^{-4} offered stable convergence without noticeable oscillation, and the Adam optimizer provided reliable gradient updates throughout the process. Because the dataset is severely imbalanced, we constructed a balanced subset of 10,000 images, evenly split between ship-containing and empty images. This reduced the risk that the network would learn a trivial "always empty" solution. During training, we monitored both binary cross-entropy and a dice-based metric to better capture segmentation quality.

The Mask R-CNN experiment required a separate training loop due to architectural differences. We fine-tuned a COCO-pretrained model using stochastic gradient descent with momentum and a scheduled learning-rate decay. Despite correct implementation and several iterations, the model converged to a poor solution, suggesting that much longer training, anchor-size adjustments or larger batch sizes might be needed.

VI. EVALUATION AND RESULTS

Model performance was assessed using the F2 score, the metric used in the original Airbus competition. U-Net achieved a significantly higher F2 score than any of the alternative approaches, reaching 0.6 on the validation set. Visual inspection also confirmed that the model successfully captured the shapes and locations of ships, even when multiple vessels appeared in the same frame.

The CNN classifier provided robust predictions on image-level ship presence, aiding error analysis and serving as a potential first-stage filter for computational efficiency.

In contrast, Mask R-CNN produced very weak outputs, with a mean F2 score of only 0.08. Predictions often missed ships entirely or segmented irrelevant regions. While this does not

invalidate the architecture itself, it suggests that much heavier training and parameter tuning would be required to reach competitive performance on this dataset.

Experiments using SAM produced similarly weak results. Although SAM occasionally generated visually plausible segments, its IoU compared to ground truth varied widely, indicating that foundation models still require adaptation for specialized satellite imagery tasks. We did not continue on this path as it requires significant computational resources [11].

VII. FUTURE PLANS AND CONCLUSION

Our experiments confirm that U-Net remains a highly effective model for maritime segmentation, especially when paired with careful data cleaning and balanced sampling strategies. The poor performance of Mask R-CNN highlights the difficulty of instance segmentation in this dataset, where object sizes vary dramatically and annotation noise is non-negligible. Additional work—particularly anchor tuning, stronger regularization, more epochs and potential tiling of high-resolution images—would be required before such an approach becomes competitive.

In future work, the U-Net could be enhanced using attention mechanisms, residual blocks or ensemble methods, and using a larger network to extract features from higher resolution images. Training on tiled patches rather than full images may help recover smaller ships. Some postprocessing should also be added to isolate separate ship masks from the unified masks.

VIII. BIBLIOGRAPHY

- [1] K. He, G. Gkioxari, P. Dollár, és R. Girshick, „Mask R-CNN”, 2018. január 24., arXiv: arXiv:1703.06870. doi: 10.48550/arXiv.1703.06870.
- [2] A. Kirillov és mtsai., „Segment Anything”, 2023. április 5., arXiv: arXiv:2304.02643. doi: 10.48550/arXiv.2304.02643.
- [3] O. Ronneberger, P. Fischer, és T. Brox, „U-Net: Convolutional Networks for Biomedical Image Segmentation”, in Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015, N. Navab, J. Hornegger, W. M. Wells, és A. F. Frangi, Szerk., Cham: Springer International Publishing, 2015, o. 234–241. doi: 10.1007/978-3-319-24574-4_28.
- [4] P. Christen, D. J. Hand, és N. Kirielle, „A Review of the F-Measure: Its History, Properties, Criticism, and Alternatives”, ACM Comput Surv, köt. 56, sz. 3, o. 73:1-73:24, okt. 2023, doi: 10.1145/3606367.
- [5] J. Long, E. Shelhamer, és T. Darrell, „Fully Convolutional Networks for Semantic Segmentation”, előadás Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2015, o. 3431–3440. Elérés: 2025. december 7. [Online]. Elérhető: https://openaccess.thecvf.com/content_cvpr_2015/html/Long_Fully_Convolutional_Networks_2015_CVPR_paper.html
- [6] X. Nie, M. Duan, H. Ding, B. Hu, és E. K. Wong, „Attention Mask R-CNN for Ship Detection and Segmentation From Remote Sensing Images”, IEEE Access, köt. 8, o. 9325–9334, 2020, doi: 10.1109/ACCESS.2020.2964540.

> REPLACE THIS LINE WITH YOUR MANUSCRIPT ID NUMBER (DOUBLE-CLICK HERE TO EDIT) <

- [7]. Usha Ruby Dr.A, „Binary cross entropy with deep learning technique for Image classification”, Int. J. Adv. Trends Comput. Sci. Eng., köt. 9, sz. 4, o. 5393–5397, aug. 2020, doi: 10.30534/ijatcse/2020/175942020.
- [8] R. Zhao és mtsai., „Rethinking Dice Loss for Medical Image Segmentation”, in 2020 IEEE International Conference on Data Mining (ICDM), nov. 2020, o. 851–860. doi: 10.1109/ICDM50108.2020.00094.
- [9] H. Rezatofighi, N. Tsoi, J. Gwak, A. Sadeghian, I. Reid, és S. Savarese, „Generalized Intersection Over Union: A Metric and a Loss for Bounding Box Regression”, előadás Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2019, o. 658–666. Elérés: 2025. december 7. [Online]. Elérhető: https://openaccess.thecvf.com/content_CVPR_2019/html/Rezatofighi_Generalized_Intersection_Over_Union_A_Metric_and_a_Loss_for_CVPR_2019_paper.html
- [10] K. Ahn, Z. Zhang, Y. Kook, és Y. Dai, „Understanding Adam Optimizer via Online Learning of Updates: Adam is FTRL in Disguise”, 2024. május 30., arXiv: arXiv:2402.01567. doi: 10.48550/arXiv.2402.01567.
- [11] „Generalized SAM: Efficient Fine-Tuning of SAM for Variable Input Image Sizes”. Elérés: 2025. december 7. [Online]. Elérhető: <https://arxiv.org/html/2408.12406v1>